

## 6.1 Assignment

October 11, 2020

### 0.0.1 Week6: Computer Vision

### 0.0.2 6.1 Assignment

Using section 5.1 in Deep Learning with Python as a guide (listing 5.3 in particular), create a ConvNet model that classifies images in the MNIST digit dataset. Save the model, predictions, metrics, and validation plots in the `dsc650/assignments/assignment06/results` directory.

### 0.0.3 5.1 - Introduction to convnets

This notebook contains the code sample found in Chapter 5, Section 1 of Deep Learning with Python. Note that the original text features far more content, in particular further explanations and figures: in this notebook, you will only find source code and related comments.

First, let's take a practical look at a very simple convnet example. We will use our convnet to classify MNIST digits, a task that you've already been through in Chapter 2, using a densely-connected network (our test accuracy then was 97.8%). Even though our convnet will be very basic, its accuracy will still blow out of the water that of the densely-connected model from Chapter 2.

The 6 lines of code below show you what a basic convnet looks like. It's a stack of Conv2D and MaxPooling2D layers. We'll see in a minute what they do concretely. Importantly, a convnet takes as input tensors of shape (image\_height, image\_width, image\_channels) (not including the batch dimension). In our case, we will configure our convnet to process inputs of size (28, 28, 1), which is the format of MNIST images. We do this via passing the argument `input_shape=(28, 28, 1)` to our first layer.

```
[1]: import keras
      keras.__version__
```

```
[1]: '2.4.3'
```

```
[2]: from keras import layers
      from keras import models

      model = models.Sequential()
      model.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
      model.add(layers.MaxPooling2D((2, 2)))
      model.add(layers.Conv2D(64, (3, 3), activation='relu'))
      model.add(layers.MaxPooling2D((2, 2)))
      model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

Let's display the architecture of our convnet so far:

```
[3]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928

Total params: 55,744  
Trainable params: 55,744  
Non-trainable params: 0

You can see above that the output of every Conv2D and MaxPooling2D layer is a 3D tensor of shape (height, width, channels). The width and height dimensions tend to shrink as we go deeper in the network. The number of channels is controlled by the first argument passed to the Conv2D layers (e.g. 32 or 64).

The next step would be to feed our last output tensor (of shape (3, 3, 64)) into a densely-connected classifier network like those you are already familiar with: a stack of Dense layers. These classifiers process vectors, which are 1D, whereas our current output is a 3D tensor. So first, we will have to flatten our 3D outputs to 1D, and then add a few Dense layers on top:

```
[4]: model.add(layers.Flatten())  
model.add(layers.Dense(64, activation='relu'))  
model.add(layers.Dense(10, activation='softmax'))
```

We are going to do 10-way classification, so we use a final layer with 10 outputs and a softmax activation. Now here's what our network looks like:

```
[5]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0

conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
-----		
max_pooling2d_1 (MaxPooling2)	(None, 5, 5, 64)	0
-----		
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
-----		
flatten (Flatten)	(None, 576)	0
-----		
dense (Dense)	(None, 64)	36928
-----		
dense_1 (Dense)	(None, 10)	650
=====		
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

As you can see, our (3, 3, 64) outputs were flattened into vectors of shape (576,), before going through two Dense layers.

Now, let's train our convnet on the MNIST digits. We will reuse a lot of the code we have already covered in the MNIST example from Chapter 2.

```
[6]: from keras.datasets import mnist
      from keras.utils import to_categorical

      (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

      train_images = train_images.reshape((60000, 28, 28, 1))
      train_images = train_images.astype('float32') / 255

      test_images = test_images.reshape((10000, 28, 28, 1))
      test_images = test_images.astype('float32') / 255

      train_labels = to_categorical(train_labels)
      test_labels = to_categorical(test_labels)
```

```
[7]: model.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
      model.fit(train_images, train_labels, epochs=5, batch_size=64)
```

```
Epoch 1/5
938/938 [=====] - 13s 14ms/step - loss: 0.1789 -
accuracy: 0.9431
Epoch 2/5
938/938 [=====] - 13s 14ms/step - loss: 0.0470 -
accuracy: 0.9857
Epoch 3/5
```

```
938/938 [=====] - 13s 14ms/step - loss: 0.0321 -  
accuracy: 0.9900  
Epoch 4/5  
938/938 [=====] - 13s 13ms/step - loss: 0.0238 -  
accuracy: 0.9923  
Epoch 5/5  
938/938 [=====] - 12s 13ms/step - loss: 0.0196 -  
accuracy: 0.9942
```

```
[7]: <tensorflow.python.keras.callbacks.History at 0x7fddf5a08dc0>
```

Let's evaluate the model on the test data:

```
[8]: test_loss, test_acc = model.evaluate(test_images, test_labels)
```

```
313/313 [=====] - 1s 4ms/step - loss: 0.0254 -  
accuracy: 0.9917
```

```
[9]: test_acc
```

```
[9]: 0.9916999936103821
```

While our densely-connected network from Chapter 2 had a test accuracy of 97.8%, our basic convnet has a test accuracy of 99.3%: we decreased our error rate by 68% (relative). Not bad!

End

```
[ ]:
```